



Sudipta Dasmohapatra  
[sd345@duke.edu](mailto:sd345@duke.edu)


# Neural Network and Deep Learning Workshop

## Spring 2020

# Example: The MNIST Fashion Database

Source: <https://github.com/zalando-research/fashion-mnist>

- The MNIST fashion dataset contains 70000 grayscale image in 10 categories
- The images show individual articles of clothing at low resolution (28x 28 pixels)
- Goal: Your goal in this project is to train a deep learning network (CNN) to classify these images of clothing



# TensorFlow

---

- Deep Learning Library: Open source, Python, Google
- 117k + users
- TensorFlow.org: Blogs, YouTube discussions, DevSummit
- Entire ecosystem of tooling to work with Keras (API), TensorFlow.js (browser); Colaboratory (cloud), TensorBoard (visualization); TensorFlow Hub (graph-modules); TensorFlow Lite (phone)
- Competitors: PyTorch, MXNet

# Example: Classify Images of Clothing

---

- The TensorFlow tutorials are written as Jupyter notebooks and run directly in Google Colab—a hosted notebook environment that requires no setup. Click the *Run in Google Colab* button.
- We will run the code for Deep Learning using Tensorflow using Google Collaborate (this provides a platform for you to run the code snippets and observe your output on google)

<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/quickstart/beginner.ipynb#scrollTo=-Z-zA3EIfKpy>

# CNN: Tensorflow Code on Google Collaborate

---

```
[3] from __future__ import absolute_import, division, print_function

# TensorFlow and tf.keras
import tensorflow as tf
from tensorflow import keras

# Helper libraries
import numpy as np
import matplotlib.pyplot as plt

print(tf.__version__)
```



1.13.1

This example uses `tf.keras`, a high-level API to build and train models in TensorFlow. The first function imports the Tensorflow and `tf.keras` and all the helper libraries for python

# CNN: Tensorflow Code on Google Collaborate



The screenshot shows the Google Colab interface for a notebook named 'beginner.ipynb'. The left sidebar displays a file explorer with a 'sample\_data' folder containing the following files: README.md, anscombe.json, california\_housing\_test.csv, california\_housing\_train.csv, mnist\_test.csv, and mnist\_train\_small.csv. The main area shows the first cell of the notebook, which contains the following Python code:

```
[1]: from __future__ import absolute_import, division, print_function

# TensorFlow and tf.keras
import tensorflow as tf
from tensorflow import keras

# Helper libraries
import numpy as np
import matplotlib.pyplot as plt

print(tf.__version__)
```

The output of the cell is displayed below the code, showing the TensorFlow version: 1.13.1.

# CNN: Import the MNIST Fashion Dataset

---

Training: 60,000 images

Test: 10,000 images

You can access the Fashion MNIST directly from TensorFlow, just import and load the data from the github library:

<https://github.com/zalando-research/fashion-mnist>

```
[4] fashion_mnist = keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```



```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
32768/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26427392/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
8192/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4423680/4422102 [=====] - 0s 0us/step
```

# Labels for Images

| Label | Class       |
|-------|-------------|
| 0     | T-shirt/top |
| 1     | Trouser     |
| 2     | Pullover    |
| 3     | Dress       |
| 4     | Coat        |
| 5     | Sandal      |
| 6     | Shirt       |
| 7     | Sneaker     |
| 8     | Bag         |
| 9     | Ankle boot  |

The *labels* are an array of integers, ranging from 0 to 9. These correspond to the *class* of clothing the image represents.

```
[5] class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',  
                  'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
```



# Explore the data

---


```
[6] train_images.shape
```

 (60000, 28, 28)


```
[7] len(train_labels)
```

 60000

```
[8] train_labels
```

 array([9, 0, 0, ..., 3, 0, 5], dtype=uint8)

```
[9] test_images.shape
```


 (10000, 28, 28)

```
[10] len(test_labels)
```

 10000

# Pre-Process the Data

```
[11] plt.figure()  
     plt.imshow(train_images[0])  
     plt.colorbar()  
     plt.grid(False)  
     plt.show(.,
```



```
[12] train_images = train_images / 255.0  
     test_images = test_images / 255.0
```

```
[13] plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[train_labels[i]])
plt.show()
```



# Build the Model: Set up Layers

```
[15] model = keras.Sequential([  
    keras.layers.Flatten(input_shape=(28, 28)),  
    keras.layers.Dense(128, activation=tf.nn.relu),  
    keras.layers.Dense(10, activation=tf.nn.softmax)  
]).
```

- First Layer: flatten layer (transforms images)
- Second ReLu Layer: 128 nodes
- Third Pooled Layer: Softmax method: 10 probability scores (10 classes)

# Build the Model: Setting

---

Add:

- **LOSS FUNCTION:** Minimize loss to move model to the right direction
- **OPTIMIZER:** Measures how model is updated based on data and loss function
- **METRICS:** Monitor the training and testing steps (accuracy= fraction of images that are correctly classified)

```
[14] model.compile(optimizer='adam',  
                  loss='sparse_categorical_crossentropy',  
                  metrics=['accuracy'])
```

# Training the Model

---

Training the neural network model requires the following steps:

- Feed the training data to the model
- The model learns to associate images and labels
- Ask the model to make predictions about a test set
- Verify that the predictions match the labels from the test\_labels array

# Training the Model

Train the neural network; Indicate number of **epochs**

An **epoch** is a full iteration over samples. The number of **epochs** is how many times the algorithm is going to run. We indicated 5 iterations for the model.

You can change that – it should be at least 2 for best results.

```
[15] model.fit(train_images, train_labels, epochs=5)
```

89% accuracy on  
training data

```
Epoch 1/5  
60000/60000 [=====] - 7s 114us/sample - loss: 0.4954 - acc: 0.8263  
Epoch 2/5  
60000/60000 [=====] - 6s 106us/sample - loss: 0.3744 - acc: 0.8648  
Epoch 3/5  
60000/60000 [=====] - 6s 106us/sample - loss: 0.3355 - acc: 0.8778  
Epoch 4/5  
60000/60000 [=====] - 6s 103us/sample - loss: 0.3115 - acc: 0.8853  
Epoch 5/5  
60000/60000 [=====] - 6s 103us/sample - loss: 0.2927 - acc: 0.8913  
<tensorflow.python.keras.callbacks.History at 0x7f3f3a185dd8>
```

# Testing the Model

---

```
[16] test_loss, test_acc = model.evaluate(test_images, test_labels)
     print('Test accuracy:', test_acc)
```

87% accuracy on  
test data

```
10000/10000 [=====] - 0s 46us/sample - loss: 0.3439 - acc: 0.8751
Test accuracy: 0.8751
```



# Making Predictions

---

87% accuracy on test data

```
[17] predictions = model.predict(test_images)
```

```
[18] predictions[0]
```

```
array([6.81665333e-05, 9.80569297e-08, 4.42758761e-07, 2.44479043e-06,  
       3.70949374e-06, 1.09513134e-01, 2.64959599e-05, 5.02966978e-02,  
       1.98009261e-03, 8.38108659e-01], dtype=float32)
```

```
[19] np.argmax(predictions[0])
```

```
9
```

```
[20] test_labels[0]
```

```
9
```

# Making Predictions

We can graph this to look at the full set of 10 channels

```
[21] def plot_image(i, predictions_array, true_label, img):
    predictions_array, true_label, img = predictions_array[i], true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img, cmap=plt.cm.binary)

    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'

    plt.xlabel("{} {:2.0f}% {}".format(class_names[predicted_label],
                                      100*np.max(predictions_array),
                                      class_names[true_label]),
              color=color)

def plot_value_array(i, predictions_array, true_label):
    predictions_array, true_label = predictions_array[i], true_label[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(predictions_array)

    thisplot[predicted_label].set_color('red')
    thisplot[true_label].set_color('blue')
```

# Making Predictions

```
[22] i = 0
      plt.figure(figsize=(6,3))
      plt.subplot(1,2,1)
      plot_image(i, predictions, test_labels, test_images)
      plt.subplot(1,2,2)
      plot_value_array(i, predictions, test_labels)
      plt.show()
```



Ankle boot 84% (Ankle boot)



Image labeled 0 in test set, the ankle boot class has been predicted as the highest probability at 84%. This image was actually an ankle boot (shown in parenthesis underneath the image) so the model predicted this image correctly

# Making Predictions

```
[23] i = 12
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(i, predictions, test_labels, test_images)
plt.subplot(1,2,2)
plot_value_array(i, predictions, test_labels)
plt.show()
```



Sandal 85% (Sneaker)




Image 12 in our test set is incorrectly predicted as a sandal with 85% accuracy or probability whereas in reality it is a sneaker.

Remember that your model is only 87% accurate so there will be some errors in prediction.

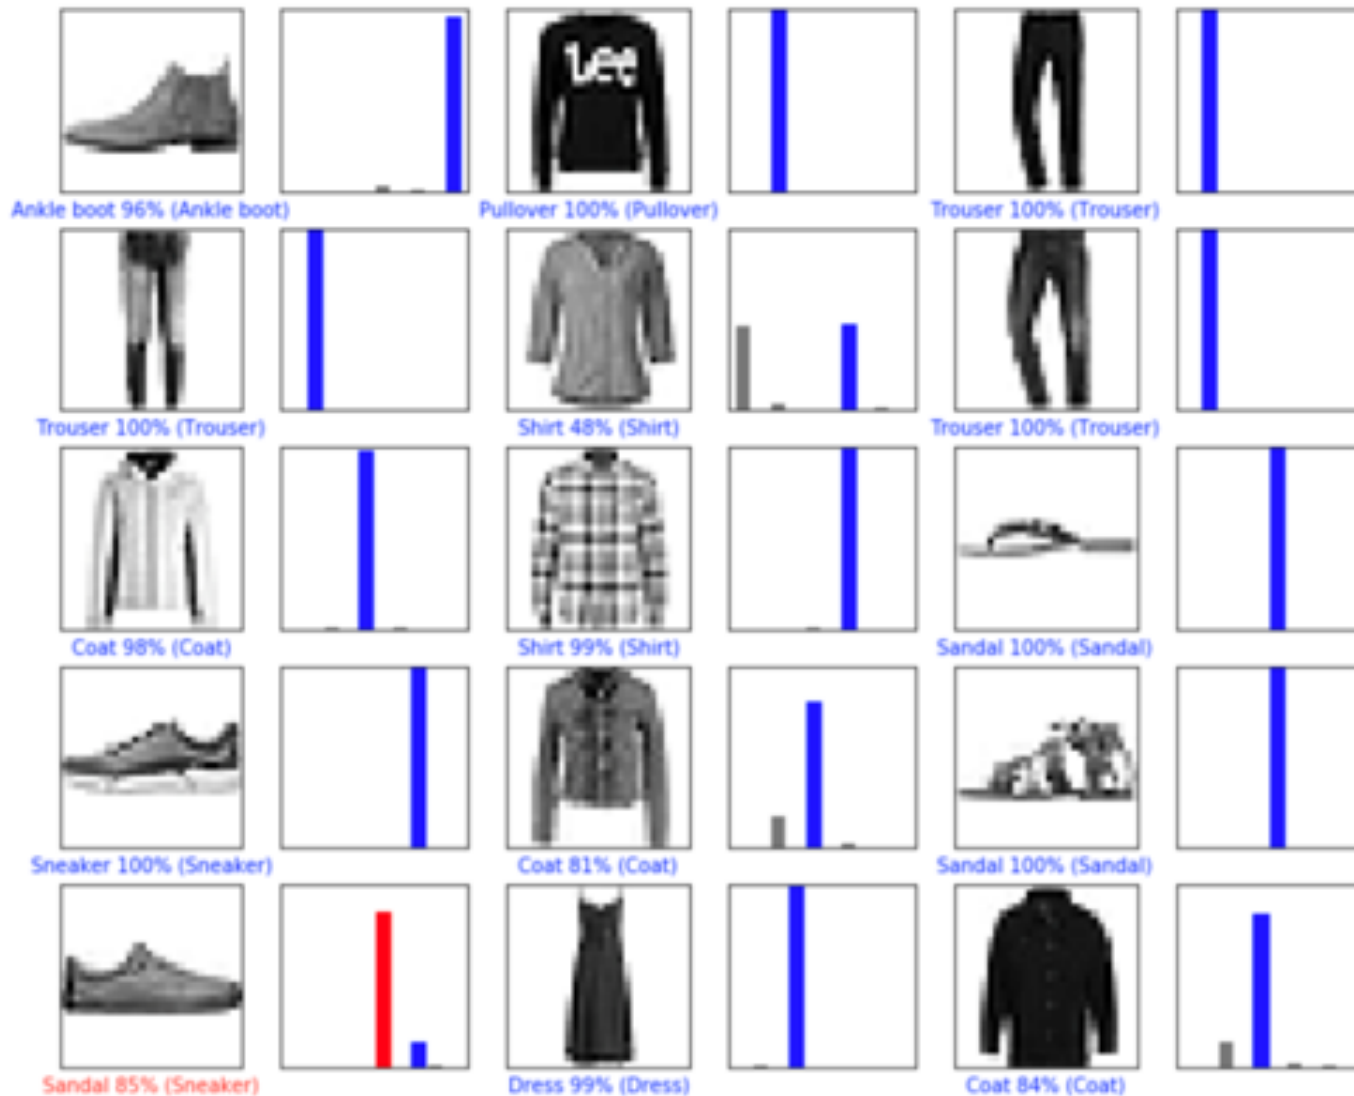
# Making Predictions

---

```
[24] # Plot the first X test images, their predicted label, and the true label
      # Color correct predictions in blue, incorrect predictions in red
      num_rows = 5
      num_cols = 3
      num_images = num_rows*num_cols
      plt.figure(figsize=(2*2*num_cols, 2*num_rows))
      for i in range(num_images):
          plt.subplot(num_rows, 2*num_cols, 2*i+1)
          plot_image(i, predictions, test_labels, test_images)
          plt.subplot(num_rows, 2*num_cols, 2*i+2)
          plot_value_array(i, predictions, test_labels)
      plt.show()
```

We ask for plotting images in 5 rows with 3 columns so total of 15 images

# Making Predictions



# Summary

---

- Deep learning concepts and principles
- Neural networks and the modeling procedure (loss functions, back propagation)
- Activation functions
- Convolutional neural networks and recurrent neural networks
- Tensor flow and example



Thanks  
[sd345@duke.edu](mailto:sd345@duke.edu)